

```
print("Welcome to Python Automation using ETL process")
```

Welcome to Python Automation using ETL process

```
import numpy as np
import pandas as pd
import matplotlib as plt
```

#Extract the Data from Flat File

```
df=pd.read_csv('TargetTransac.csv',encoding='latin1')
```

```
df.head(3)
```

	row	orderid	orderdate	ordermonth	ship_date	delay	\
0	1	CA-2016-152156	08-11-2016	Nov	11-11-2016	3	
1	2	CA-2016-152156	08-11-2016	Nov	11-11-2016	3	
2	3	CA-2016-138688	12-06-2016	Jun	16-06-2016	4	

	ship_mode	cust_id	cust_name	segment	...
category \					
0	Second Class	CG-12520	Claire Gute	Consumer	...
Furniture					
1	Second Class	CG-12520	Claire Gute	Consumer	...
Furniture					
2	Second Class	DV-13045	Darrin Van Huff	Corporate	... Office Supplies

	sub_category	product_name
sales \		
0	Bookcases	Bush Somerset Collection Bookcase
261.96		
1	Chairs	Hon Deluxe Fabric Upholstered Stacking Chairs,...
731.94		
2	Labels	Self-Adhesive Address Labels for Typewriters b...
14.62		

	quantity	discount	profit	Unnamed: 23	Unnamed: 24	Unnamed: 25
0	2.0	0.0	41.9136	NaN	NaN	NaN
1	3.0	0.0	219.5820	NaN	NaN	NaN
2	2.0	0.0	6.8714	NaN	NaN	NaN

[3 rows x 26 columns]

```
df.tail(3)
```

	row	orderid	orderdate	ordermonth	ship_date	delay	\
10207	8323	CA-2016-130778	19-11-2016	Nov	25-11-2016	6	
10208	8324	CA-2016-130778	19-11-2016	Nov	25-11-2016	6	

```
10209  8325  CA-2017-144456  08-09-2017          Sep  09-09-2017          1
```

```
          ship_mode  cust_id  cust_name
segment ... \
10207  Standard Class  ND-18370  Natalie DeCherney  Consumer ...
10208  Standard Class  ND-18370  Natalie DeCherney  Consumer ...
10209    First Class  FC-14245    Frank Carlisle  Home Office ...
```

```
          category sub_category \
10207  Office Supplies  Appliances
10208  Office Supplies      Paper
10209  Office Supplies      Storage
```

```
          product_name  sales
quantity \
10207          Disposable Triple-Filter Dust Bags  8.74
2.0
10208  Recycled Desk Saver Line "While You Were Out" ...  44.75
5.0
10209    Decoflex Hanging Personal Folder File, Blue  61.68
5.0
```

```
          discount  profit  Unnamed: 23  Unnamed: 24  Unnamed: 25
10207          0.0    2.2724          NaN          NaN          NaN
10208          0.0   20.5850          NaN          NaN          NaN
10209          0.2    5.3970          NaN          NaN          NaN
```

```
[3 rows x 26 columns]
```

```
df.shape
```

```
(10210, 26)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 10210 entries, 0 to 10209
```

```
Data columns (total 26 columns):
```

```
#    Column          Non-Null Count  Dtype
---  -
0    row            10210 non-null    int64
1    orderid         10210 non-null    object
2    orderdate       10210 non-null    object
3    ordermonth      10210 non-null    object
4    ship_date       10210 non-null    object
5    delay           10210 non-null    int64
```

```

6  ship_mode      10210 non-null object
7  cust_id        10210 non-null object
8  cust_name      10210 non-null object
9  segment        10210 non-null object
10 country        10210 non-null object
11 city           10210 non-null object
12 state          10210 non-null object
13 postal_code    10210 non-null int64
14 region         10210 non-null object
15 product_id     10210 non-null object
16 category       10210 non-null object
17 sub_category   10210 non-null object
18 product_name   10210 non-null object
19 sales          10200 non-null float64
20 quantity       10200 non-null float64
21 discount       10199 non-null float64
22 profit         10194 non-null float64
23 Unnamed: 23    0 non-null float64
24 Unnamed: 24    0 non-null float64
25 Unnamed: 25    1 non-null float64
dtypes: float64(7), int64(3), object(16)
memory usage: 2.0+ MB

```

```
df.columns
```

```

Index(['row', 'orderid', 'orderdate', 'ordermonth', 'ship_date',
      'delay',
      'ship_mode', 'cust_id', 'cust_name', 'segment', 'country',
      'city',
      'state', 'postal_code', 'region', 'product_id', 'category',
      'sub_category', 'product_name', 'sales', 'quantity',
      'discount',
      'profit', 'Unnamed: 23', 'Unnamed: 24', 'Unnamed: 25'],
      dtype='object')

```

#Transform the Dataset || Data Preprocessing

```

df.duplicated().sum()
216
df.drop_duplicates(inplace=True)
df.shape
(9994, 26)
df.isnull().sum()
row      0
orderid  0

```

```

orderdate      0
ordermonth     0
ship_date      0
delay          0
ship_mode      0
cust_id        0
cust_name      0
segment        0
country        0
city          0
state          0
postal_code    0
region         0
product_id     0
category       0
sub_category   0
product_name   0
sales          10
quantity       10
discount       11
profit         16
Unnamed: 23    9994
Unnamed: 24    9994
Unnamed: 25    9993
dtype: int64

```

```

df.drop(['Unnamed: 23', 'Unnamed: 24', 'Unnamed: 25'], axis=1,
inplace=True)

```

```

df.columns

```

```

Index(['row', 'orderid', 'orderdate', 'ordermonth', 'ship_date',
      'delay',
      'ship_mode', 'cust_id', 'cust_name', 'segment', 'country',
      'city',
      'state', 'postal_code', 'region', 'product_id', 'category',
      'sub_category', 'product_name', 'sales', 'quantity',
      'discount',
      'profit'],
      dtype='object')

```

```

df.fillna({
    'sales': df['sales'].mean(),
    'quantity': df['quantity'].median(),
    'discount': df['discount'].mean(),
    'profit': df['profit'].median()
}, inplace=True)

```

```

df.isnull().sum()

```

```
row          0
orderid      0
orderdate    0
ordermonth   0
ship_date    0
delay        0
ship_mode    0
cust_id      0
cust_name    0
segment      0
country      0
city         0
state        0
postal_code  0
region       0
product_id   0
category     0
sub_category 0
product_name 0
sales        0
quantity     0
discount     0
profit       0
dtype: int64
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 9994 entries, 0 to 9993
```

```
Data columns (total 23 columns):
```

#	Column	Non-Null Count	Dtype
0	row	9994 non-null	int64
1	orderid	9994 non-null	object
2	orderdate	9994 non-null	object
3	ordermonth	9994 non-null	object
4	ship_date	9994 non-null	object
5	delay	9994 non-null	int64
6	ship_mode	9994 non-null	object
7	cust_id	9994 non-null	object
8	cust_name	9994 non-null	object
9	segment	9994 non-null	object
10	country	9994 non-null	object
11	city	9994 non-null	object
12	state	9994 non-null	object
13	postal_code	9994 non-null	int64
14	region	9994 non-null	object
15	product_id	9994 non-null	object
16	category	9994 non-null	object
17	sub_category	9994 non-null	object

```

18 product_name 9994 non-null object
19 sales        9994 non-null float64
20 quantity    9994 non-null float64
21 discount     9994 non-null float64
22 profit       9994 non-null float64
dtypes: float64(4), int64(3), object(16)
memory usage: 1.8+ MB

```

```

df.rename(columns={'ship_date':'shipdate',
'ship_mode':'shipmode','cust_id':'customerid','cust_name':'customernam
e','postal_code':'postalcode', 'product_id':'productid',
'sub_category': 'subcategory', 'product_name':'productname'},inplace=
True)

```

```
df.columns
```

```

Index(['row', 'orderid', 'orderdate', 'ordermonth', 'shipdate',
'delay',
      'shipmode', 'customerid', 'customername', 'segment', 'country',
'city',
      'state', 'postalcode', 'region', 'productid', 'category',
'subcategory',
      'productname', 'sales', 'quantity', 'discount', 'profit'],
      dtype='object')

```

```
data=(df[['profit','sales','discount']]).round(1)
```

```
print(data)
```

	profit	sales	discount
0	41.9	262.0	0.0
1	219.6	731.9	0.0
2	6.9	14.6	0.0
3	-383.0	957.6	0.4
4	2.5	22.4	0.2
...
9989	4.1	25.2	0.2
9990	15.6	92.0	0.0
9991	19.4	258.6	0.2
9992	13.3	29.6	0.0
9993	72.9	243.2	0.0

```
[9994 rows x 3 columns]
```

```
df.drop(['profit', 'sales', 'discount'], axis=1, inplace=True)
```

```
df.columns
```

```

Index(['row', 'orderid', 'orderdate', 'ordermonth', 'shipdate',
'delay',
      'shipmode', 'customerid', 'customername', 'segment', 'country',

```

```

'city',
    'state', 'postalcode', 'region', 'productid', 'category',
    'subcategory',
    'productname', 'quantity'],
    dtype='object')

data_new = pd.concat([df, data], axis=1)

data_new.columns

Index(['row', 'orderid', 'orderdate', 'ordermonth', 'shipdate',
      'delay',
      'shipmode', 'customerid', 'customername', 'segment', 'country',
      'city',
      'state', 'postalcode', 'region', 'productid', 'category',
      'subcategory',
      'productname', 'quantity', 'profit', 'sales', 'discount'],
      dtype='object')

print(data_new[['profit', 'sales', 'discount']])

```

	profit	sales	discount
0	41.9	262.0	0.0
1	219.6	731.9	0.0
2	6.9	14.6	0.0
3	-383.0	957.6	0.4
4	2.5	22.4	0.2
...
9989	4.1	25.2	0.2
9990	15.6	92.0	0.0
9991	19.4	258.6	0.2
9992	13.3	29.6	0.0
9993	72.9	243.2	0.0

[9994 rows x 3 columns]

```
data_new.shape
```

```
(9994, 23)
```

```
data_new.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Index: 9994 entries, 0 to 9993
Data columns (total 23 columns):
#   Column          Non-Null Count  Dtype
---  -
0   row             9994 non-null  int64
1   orderid         9994 non-null  object
2   orderdate       9994 non-null  object
3   ordermonth      9994 non-null  object

```

```

4  shipdate      9994 non-null  object
5  delay         9994 non-null  int64
6  shipmode      9994 non-null  object
7  customerid    9994 non-null  object
8  customername  9994 non-null  object
9  segment       9994 non-null  object
10 country       9994 non-null  object
11 city          9994 non-null  object
12 state         9994 non-null  object
13 postalcode    9994 non-null  int64
14 region        9994 non-null  object
15 productid     9994 non-null  object
16 category      9994 non-null  object
17 subcategory   9994 non-null  object
18 productname   9994 non-null  object
19 quantity      9994 non-null  float64
20 profit        9994 non-null  float64
21 sales         9994 non-null  float64
22 discount      9994 non-null  float64

```

```
dtypes: float64(4), int64(3), object(16)
```

```
memory usage: 1.8+ MB
```

```
data_new['orderdate'] = pd.to_datetime(data_new['orderdate'],
format='mixed', dayfirst=True)
```

```
data_new['shipdate'] = pd.to_datetime(data_new['shipdate'],
format='mixed', dayfirst=True)
```

```
data_new.columns == 'delay'
```

```
array([False, False, False, False, False,  True, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False, False, False])
```

```
data_new.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 9994 entries, 0 to 9993
```

```
Data columns (total 23 columns):
```

#	Column	Non-Null Count	Dtype
0	row	9994 non-null	int64
1	orderid	9994 non-null	object
2	orderdate	9994 non-null	datetime64[ns]
3	ordermonth	9994 non-null	object
4	shipdate	9994 non-null	datetime64[ns]
5	delay	9994 non-null	int64
6	shipmode	9994 non-null	object
7	customerid	9994 non-null	object
8	customername	9994 non-null	object
9	segment	9994 non-null	object


```

10  country      9994 non-null  object
11  city         9994 non-null  object
12  state        9994 non-null  object
13  postalcode   9994 non-null  int64
14  region       9994 non-null  object
15  productid    9994 non-null  object
16  category     9994 non-null  object
17  subcategory  9994 non-null  object
18  productname  9994 non-null  object
19  quantity     9994 non-null  float64
20  profit       9994 non-null  float64
21  sales        9994 non-null  float64
22  discount     9994 non-null  float64
dtypes: datetime64[ns](2), float64(4), int64(3), object(14)
memory usage: 1.8+ MB

data_new['delay']=data_new['orderdate']- data_new['shipdate']

data_new.columns
Index(['row', 'orderid', 'orderdate', 'ordermonth', 'shipdate',
      'delay',
      'shipmode', 'customerid', 'customername', 'segment', 'country',
      'city',
      'state', 'postalcode', 'region', 'productid', 'category',
      'subcategory',
      'productname', 'quantity', 'profit', 'sales', 'discount'],
      dtype='object')

filtered = data_new[data_new['profit'] < 0]
print(filtered['profit'].value_counts())

profit
-4.8      15
-1.4      15
-2.7      15
-3.0      13
-2.0      13
..
-1480.0    1
-427.4     1
-40.8      1
-46.9      1
-54.3      1
Name: count, Length: 829, dtype: int64

print(len(filtered))

1855

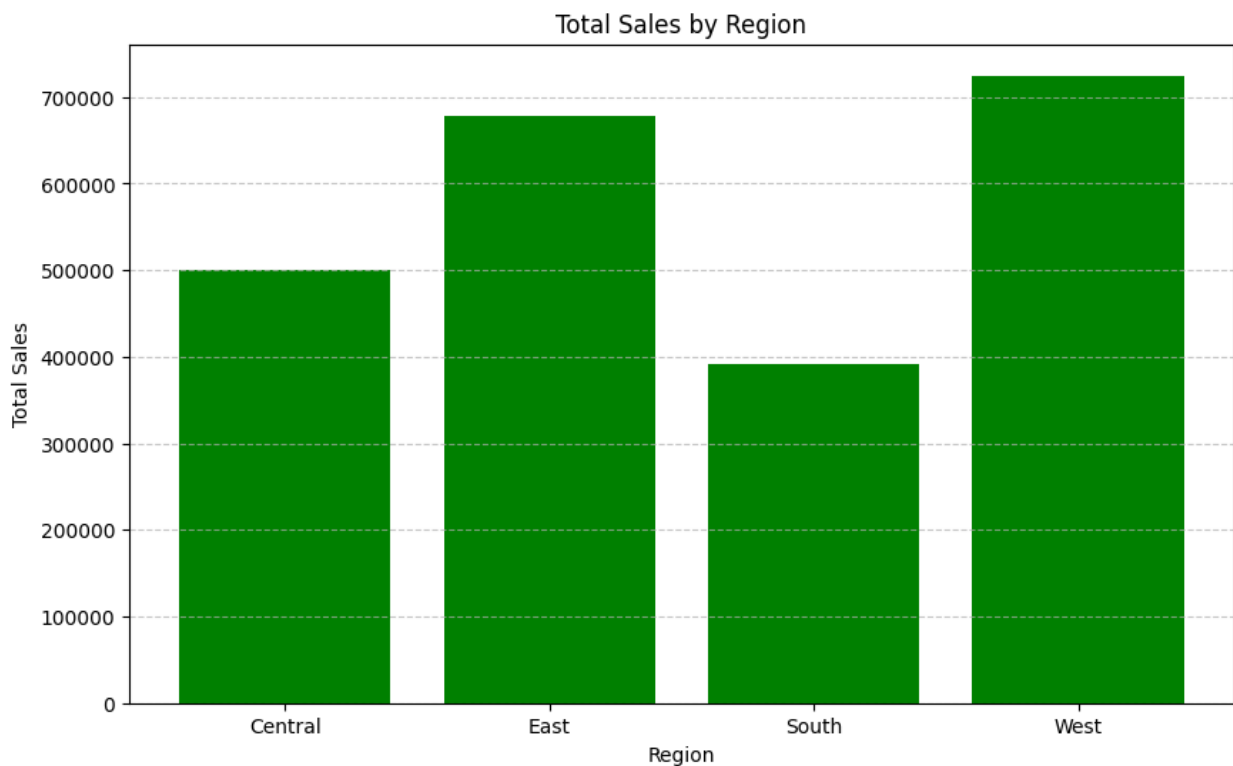
```

```
sales_summary= data_new.groupby("region")["sales"].sum().reset_index()
print(sales_summary)
```

	region	sales
0	Central	501006.8
1	East	678607.3
2	South	391728.8
3	West	724755.3

```
import matplotlib.pyplot as plt
```

```
plt.figure(figsize=(10, 6))
plt.bar(sales_summary['region'], sales_summary['sales'], color=
'green')
plt.title("Total Sales by Region")
plt.xlabel("Region")
plt.ylabel("Total Sales")
plt.grid(axis="y", linestyle= "--", alpha=0.7)
plt.show()
```

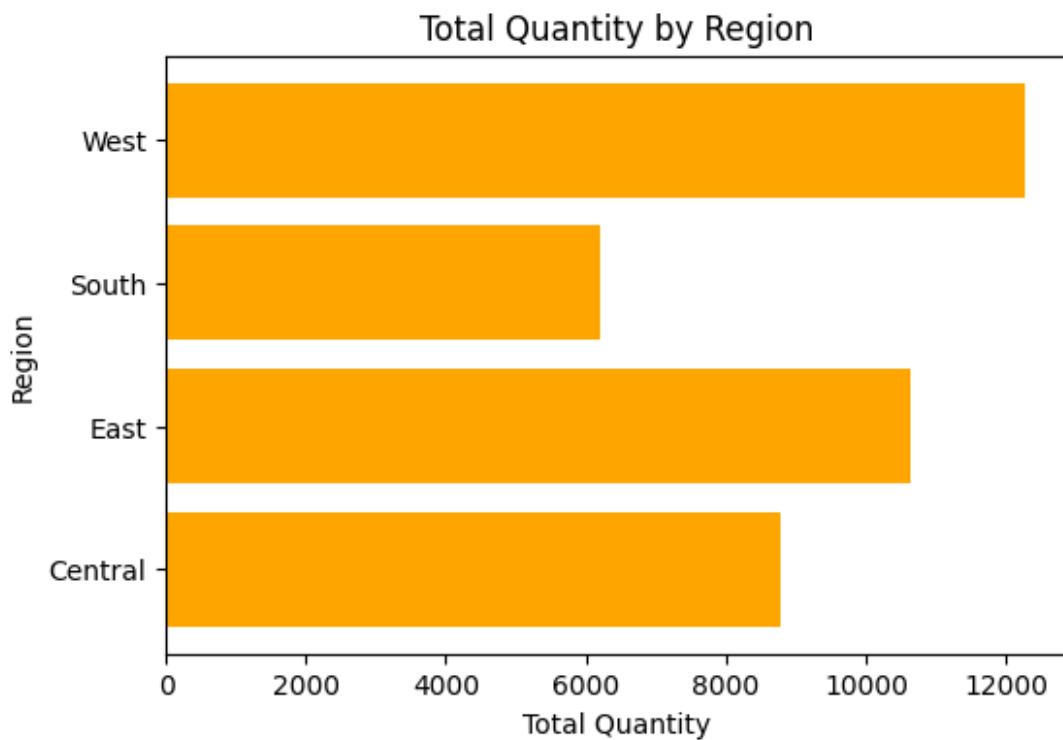


```
quantity_summary= data_new.groupby("region")
["quantity"].sum().reset_index()
print(quantity_summary)
```

	region	quantity
0	Central	8766.0

```
1    East    10622.0
2    South    6209.0
3    West    12263.0
```

```
plt.figure(figsize=(6, 4))
plt.barh(quantity_summary['region'], quantity_summary['quantity'],
color='orange')
plt.title("Total Quantity by Region")
plt.xlabel("Total Quantity")
plt.ylabel("Region")
plt.show()
```



Load the dataset for further analysis

```
data_new.to_csv('detail.csv', index=False)
```